A.I.: Informed Search Algorith

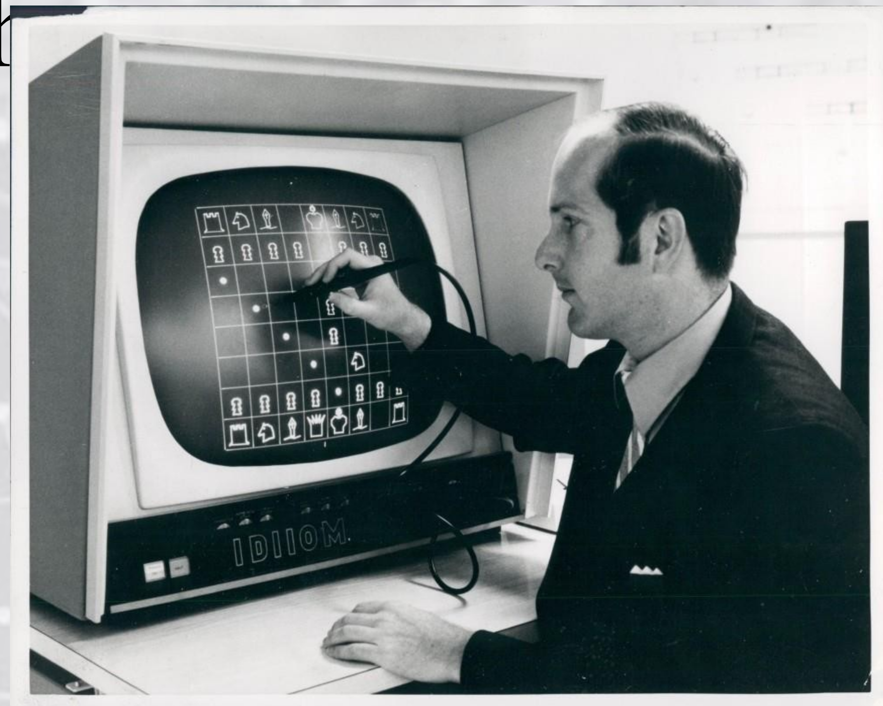

Outline

- Best-first search
- Greedy best-first search
- A^* search
- Heuristics

Overview

- **Informed Search**: uses problem-specific knowledge.

- General approach: **best-first search**; an instance of TREE-SEARCH (or GRAPH-SEARCH) – where a search strategy is defined by picking the order of node expansion.

- With best-first, node is selected for expansion based on **evaluation function** *f(n)*.

- Evaluation function is a *cost estimate*; expand lowest cost node first (same as uniform-cost search but we replace *g* with *f*).

Overview (cont'd)

- The choice of *f* determines the search strategy (one can show that best-first tree search includes DFS as a special case).

- Often, for best-first algorithms, f is defined in terms of a **heuristic function**, h(n).

  h(n) = *estimated* cost of the cheapest path from the state at node *n* to a *goal state*. (for goal state: *h(n)*=0)

- Heuristic functions are the most common form in which **additional knowledge** of the problem is passed to the search algorithm.

Overview (cont'd)

- Best-First Search algorithms constitute a large family of algorithms, with different evaluation functions.
  - Each has a heuristic function h(n)

- Example: in route planning the estimate of the cost of the cheapest path might be the straight line distance between two cities.

Recall:

- *g(n)* = cost from the initial state to the current state *n*.
- *h(n)* = estimated cost of the cheapest path from node *n* to a goal node.
- *f(n)* = evaluation function to select a node for expansion (usually the lowest cost node)

Best-First Search

- Idea: use an evaluation function *f(n)* for each node
  - f(n) provides an estimate for the total cost.
  - Expand the node n with smallest f(n).

- <u>Implementation</u>:
  Order the nodes in the frontier increasing order of cost.

- Special cases:
  - Greedy best-first search
  - A^* search

Greedy best-first search

- Evaluation function $f(n) = h(n)$ (*heuristic*), the estimate of cost from n to *goal*.
- We use the straight-line distance heuristic: $h_{SLD}(n)$ = straight-line distance from n to Bucharest.
- Note that the heuristic values cannot be computed from the problem description itself!
- In addition, we require **extrinsic knowledge** to understand that h_{SLD} is correlated with the actual road distances, making it a useful heuristic.
- Greedy best-first search expands the node that appears to be closest to goal.

Romania with step costs in km

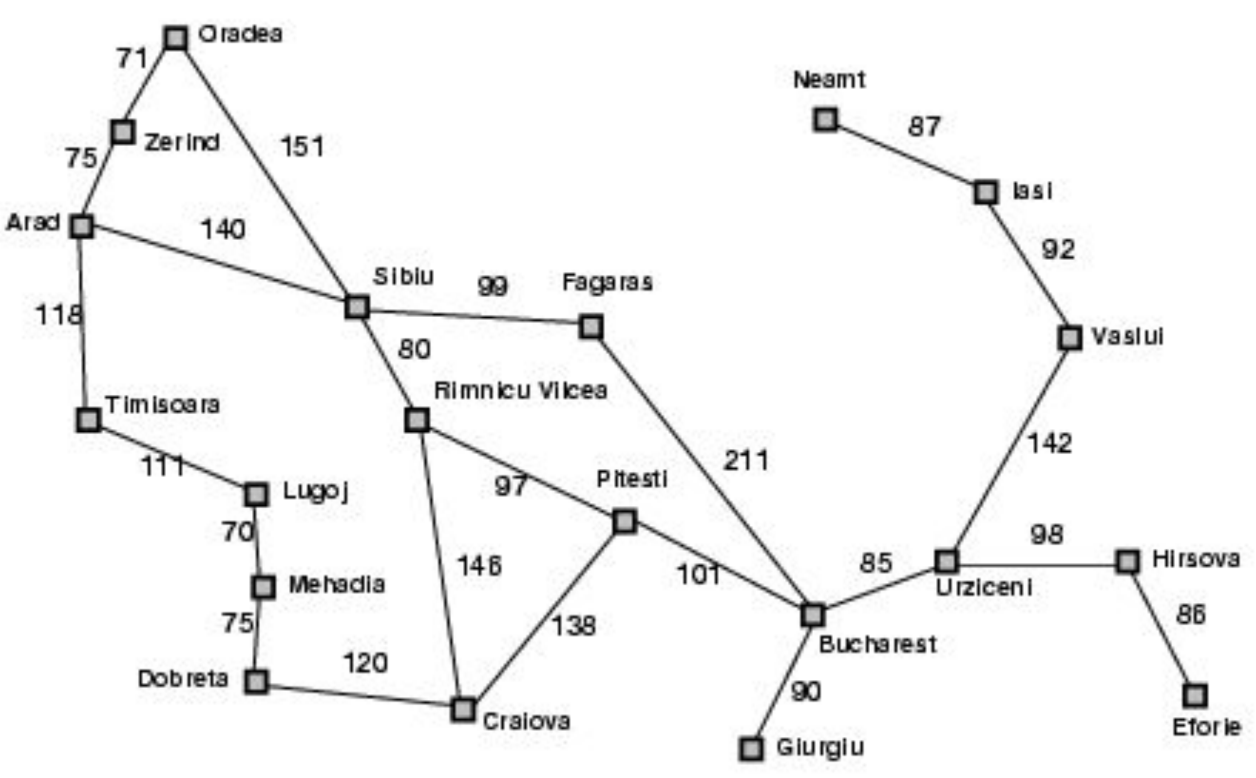

Straight–line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy best-first search exa

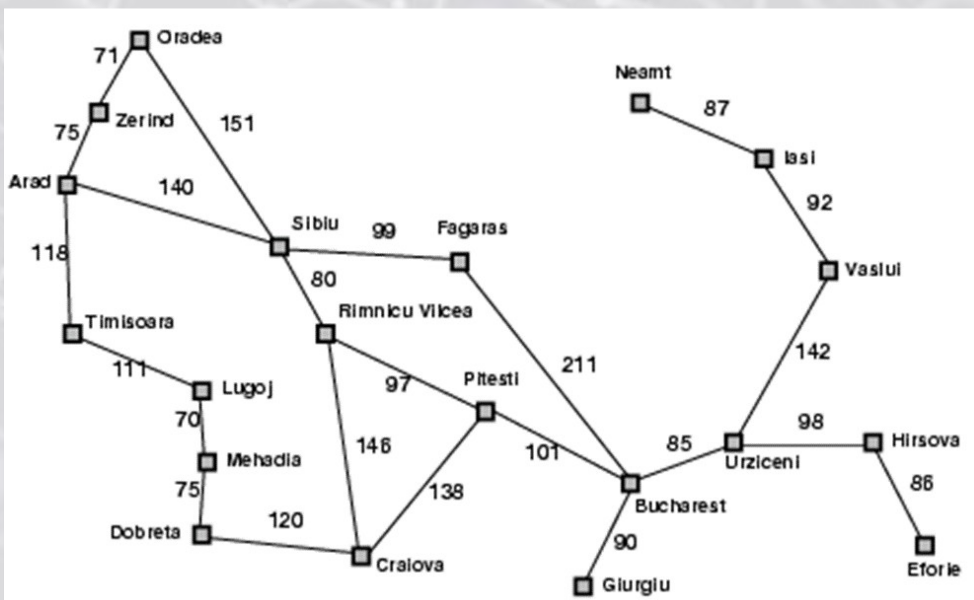

Straight–line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy best-first search example

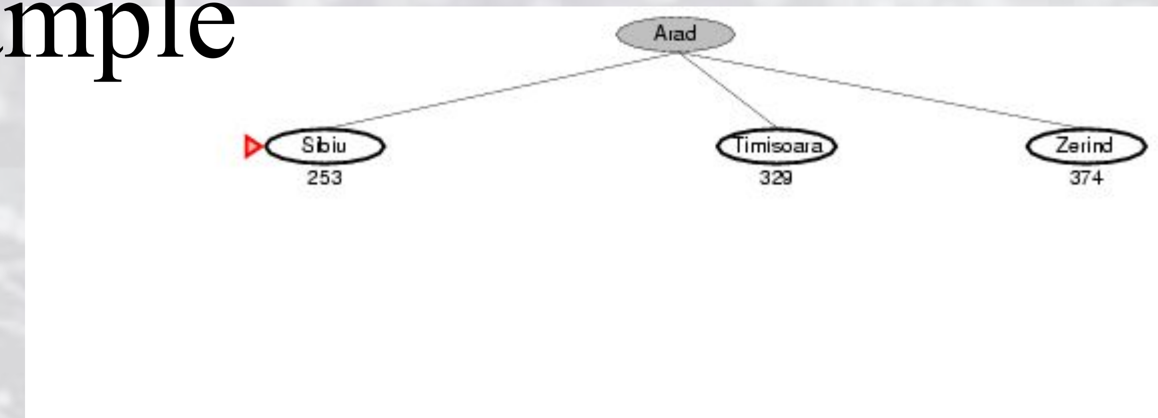

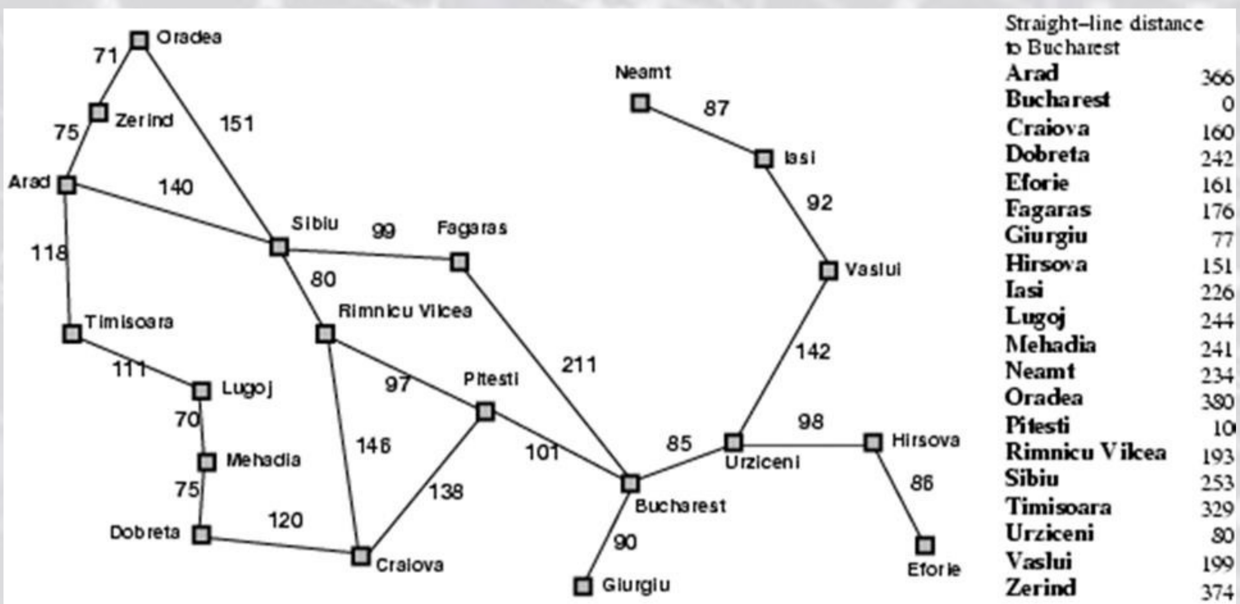

Greedy best-first search example

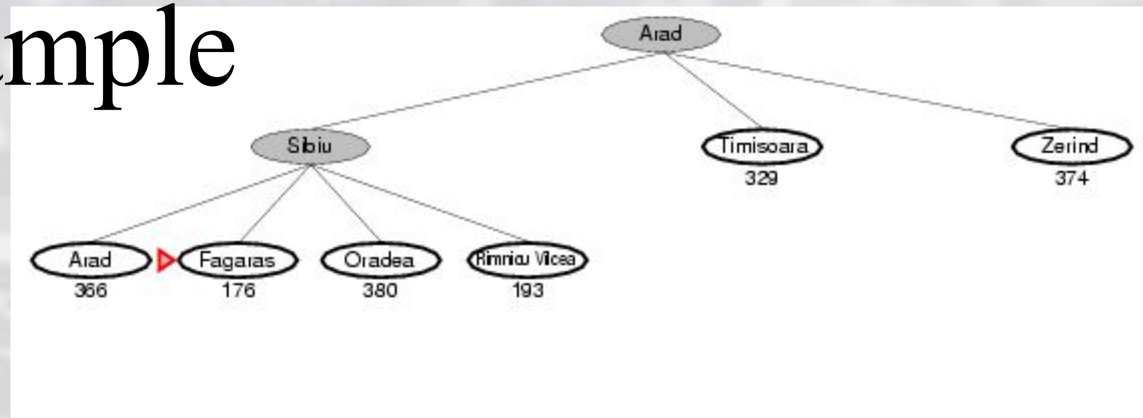

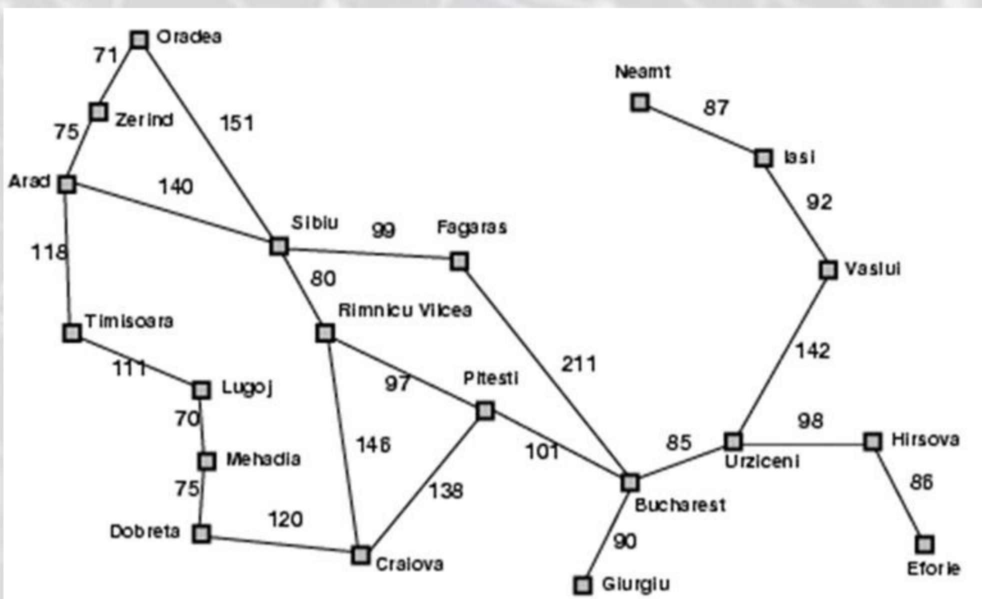

Straight–line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Greedy best-first search example

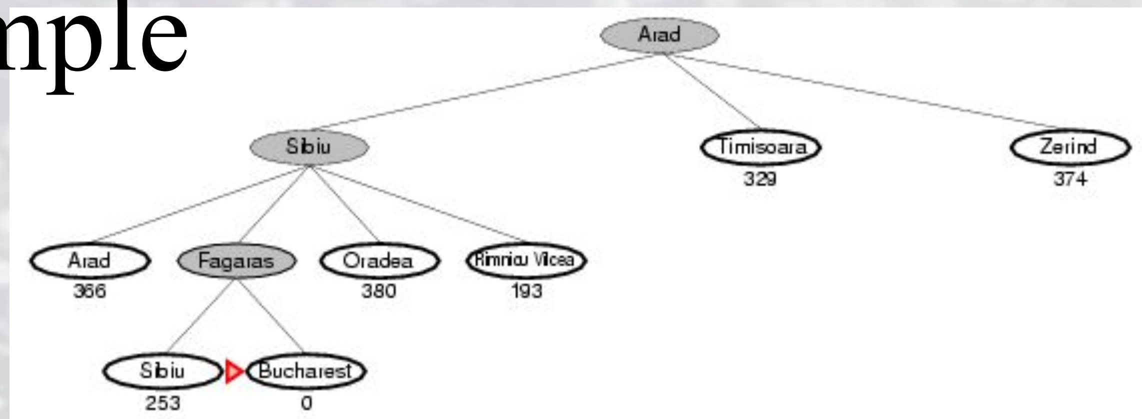

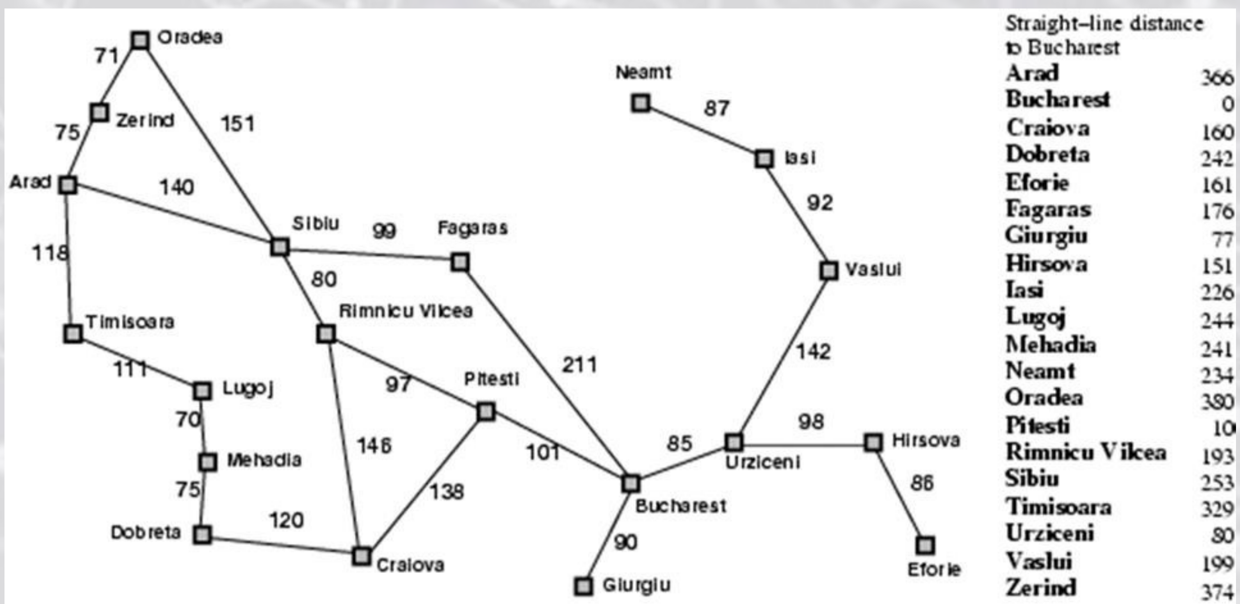

Greedy best-first

- GBFS is search *incomplete*!

- Why?

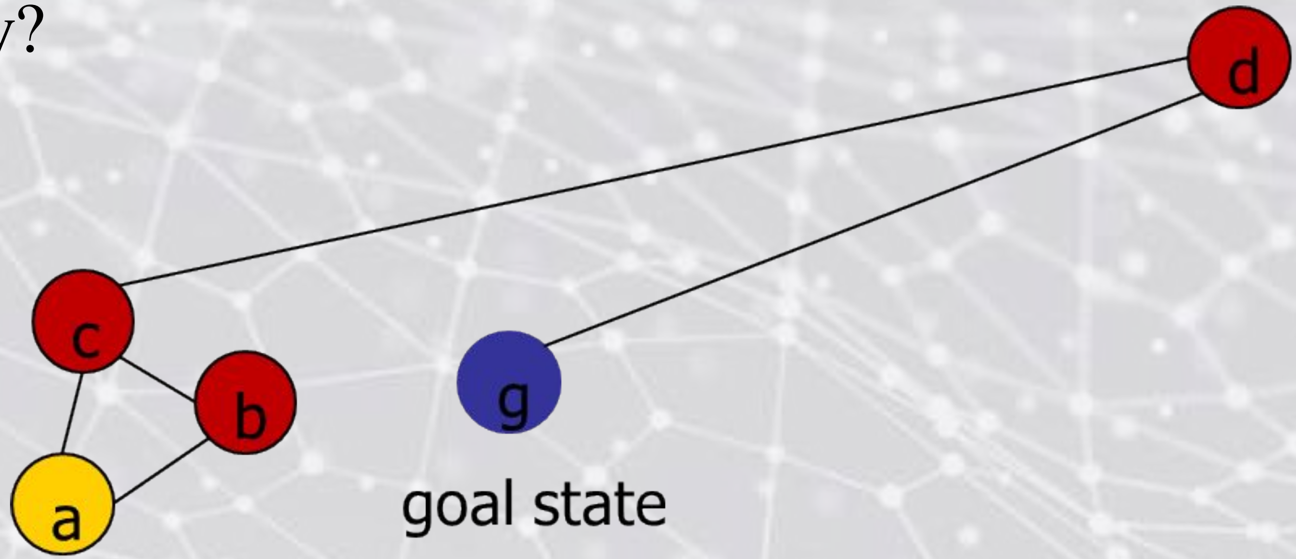

- Graph-Search version is, however, complete in *finite* spaces.

Properties of greedy best-first search

- <u>Complete?</u> No – can get stuck in loops, e.g., □ Neamt □ Iasi □ Neamt □ Iasi
- <u>Time?</u> *$O(b^m)$*, *(*in worst case*)* but a good heuristic can give dramatic improvement (*m* is max depth of search space).
- <u>Space?</u> *$O(b^m)$* -- keeps all nodes in memory.
- <u>Optimal?</u> No (not guaranteed to render lowest cost solution).

A* Search

- Most widely-known form of best-first search.

- It evaluates nodes by combining $g(n)$, the cost to reach the node, and $h(n)$, the cost to get from the node to the goal:

$f(n) = g(n) + h(n)$ (estimated cost of cheapest solution through n).

- A reasonable strategy: try node with the lowest $g(n) + h(n)$ value!

A* search

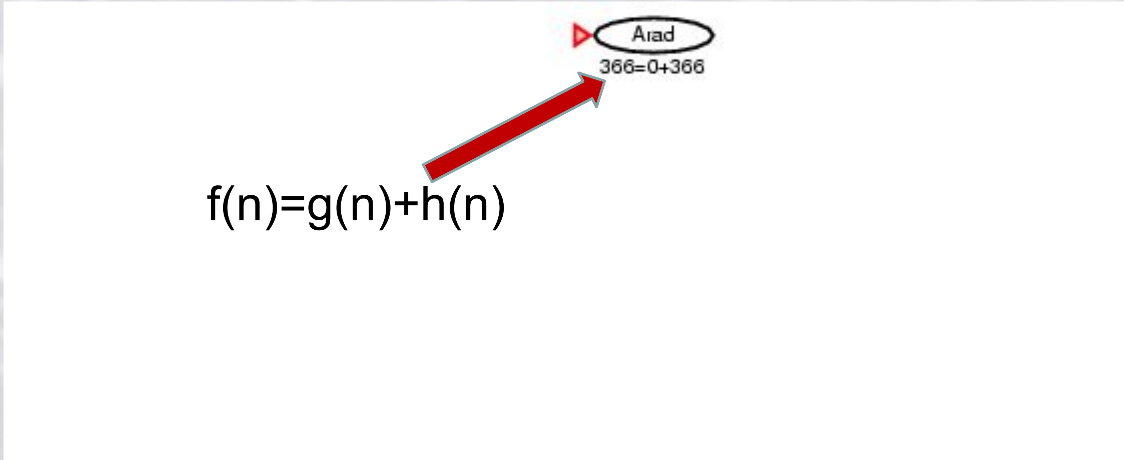

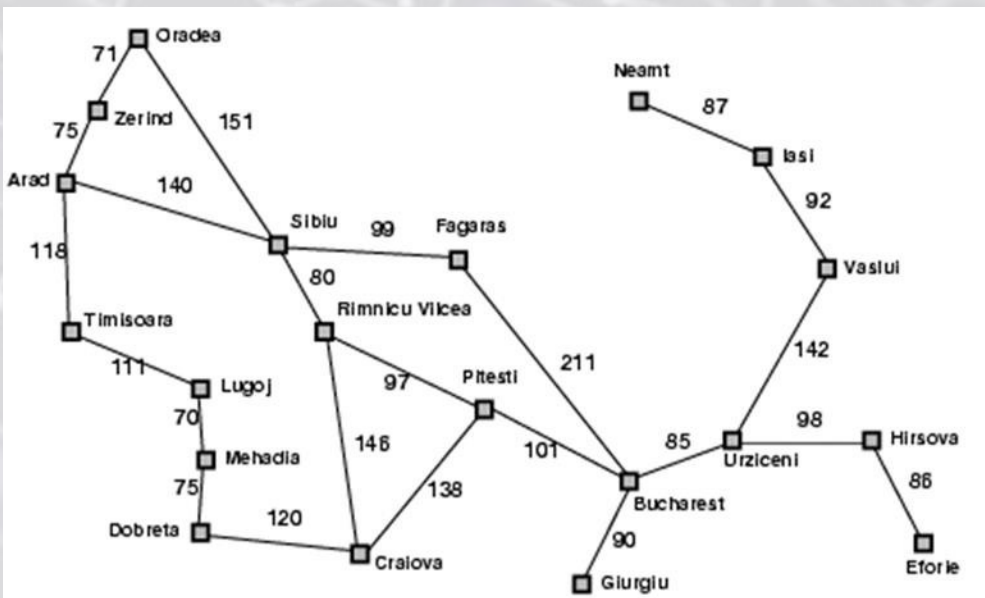

Straight–line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

A* search example

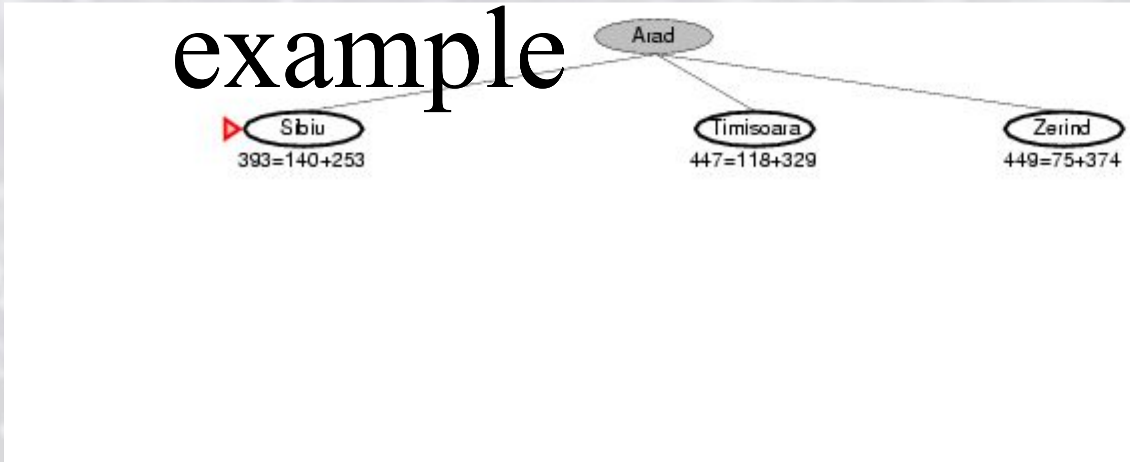

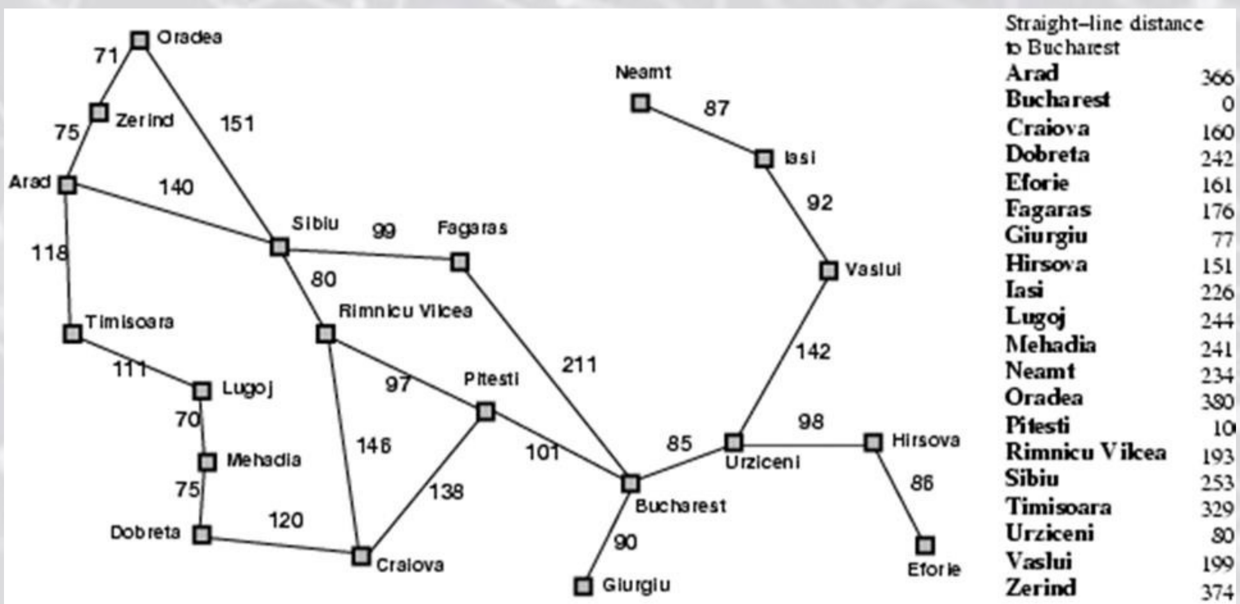

A* search

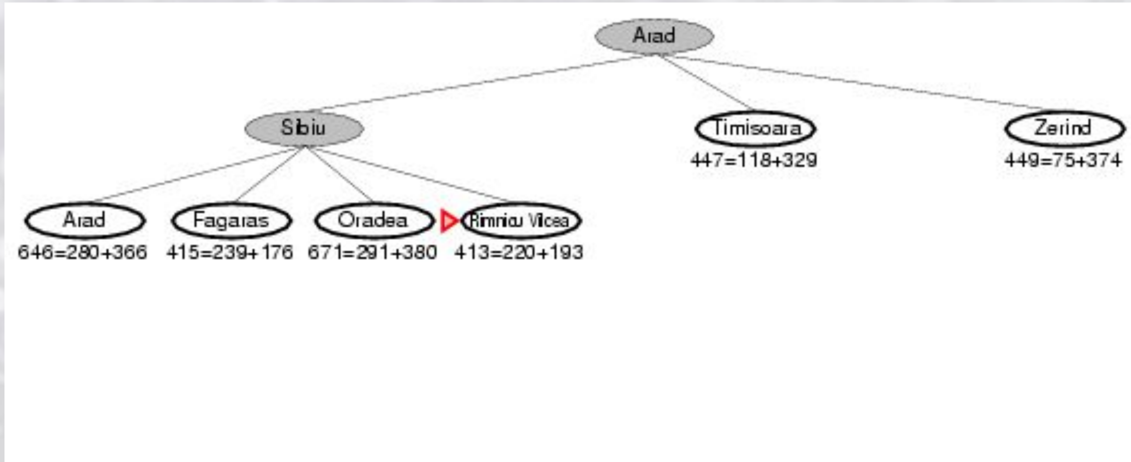

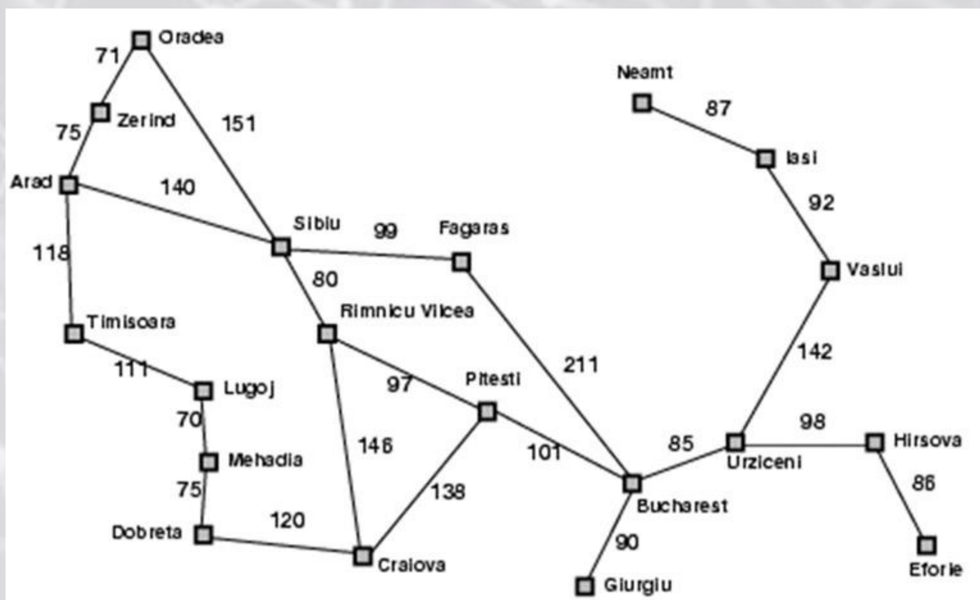

Straight–line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	10
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

A* search

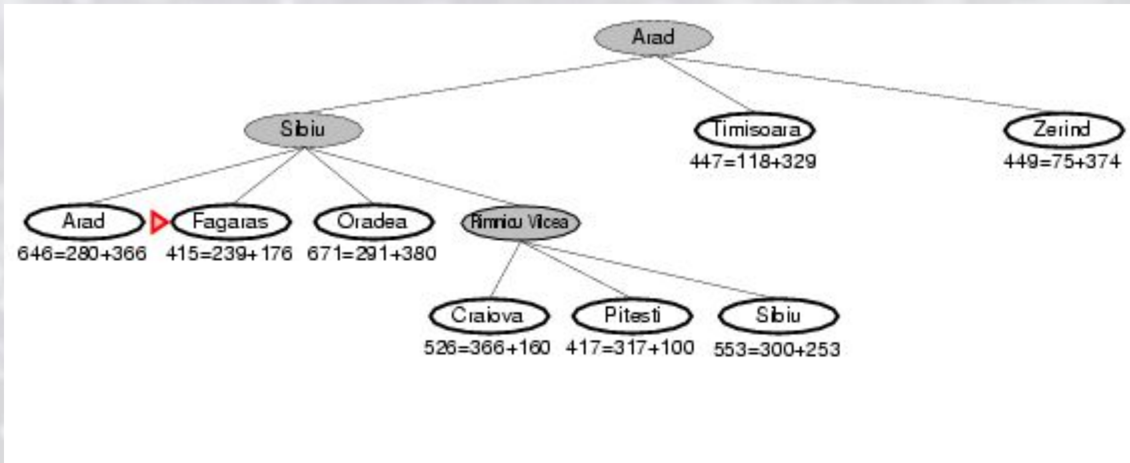
Arad
Sibiu
Timisoara
447=118+329
Zerind
449=75+374
Arad
646=280+366
Fagaras
415=239+176
Oradea
671=291+380
Rimnicu Vilcea
Craiova
526=366+160
Pitesti
417=317+100
Sibiu
553=300+253

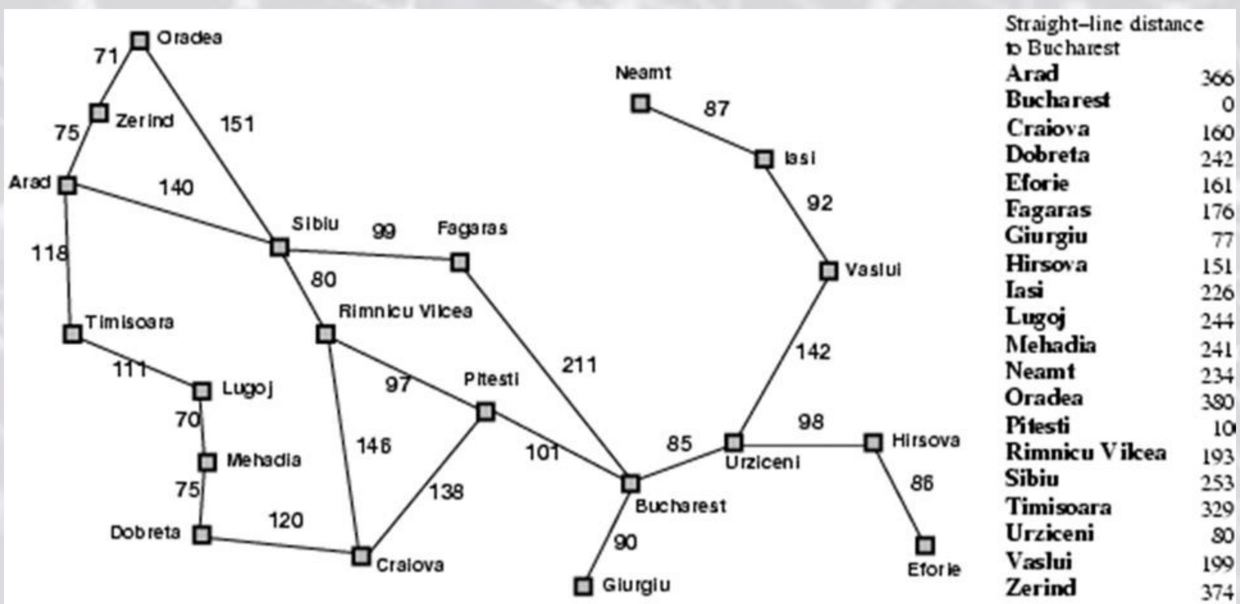
Oradea
71
Zerind
75
151
Arad
140
118
Sibiu
99
Fagaras
80
Rimnicu Vilcea
Timisoara
111
Lugoj
97
Pitesti
211
70
Mehadia
146
101
75
138
Dobreta
120
Craiova
Bucharest
90
Giurgiu
85
Urziceni
98
Hirsova
86
Eforie
142
Vaslui
92
Iasi
87
Neamt
Straight–line distance to Bucharest
Arad 366
Bucharest 0
Craiova 160
Dobreta 242
Eforie 161
Fagaras 176
Giurgiu 77
Hirsova 151
Iasi 226
Lugoj 244
Mehadia 241
Neamt 234
Oradea 380
Pitesti 10
Rimnicu Vilcea 193
Sibiu 253
Timisoara 329
Urziceni 80
Vaslui 199
Zerind 374

A* search

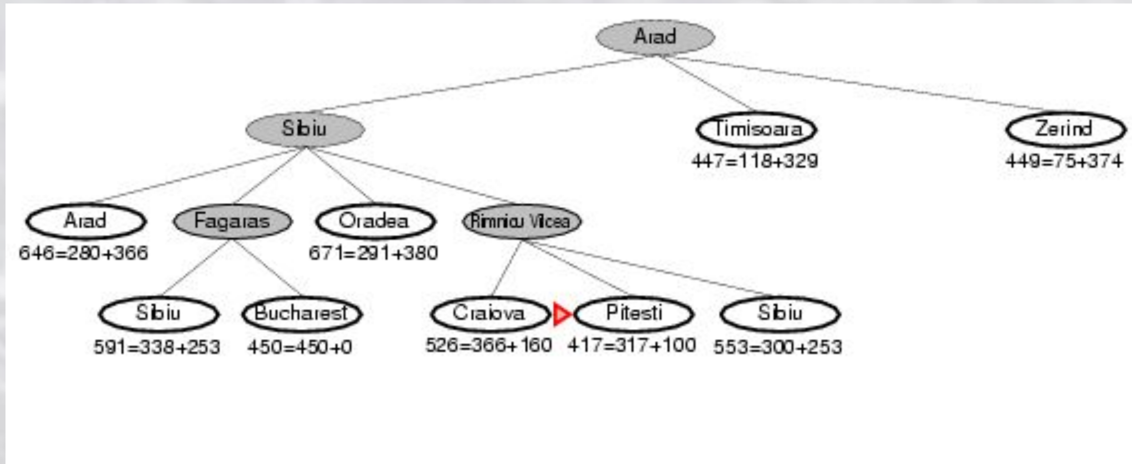

Arad
Sibiu
Timisoara
447=118+329
Zerind
449=75+374
Arad
646=280+366
Fagaras
Oradea
671=291+380
Rimnicu Vilcea
Sibiu
591=338+253
Bucharest
450=450+0
Craiova
526=366+160
Pitesti
417=317+100
Sibiu
553=300+253

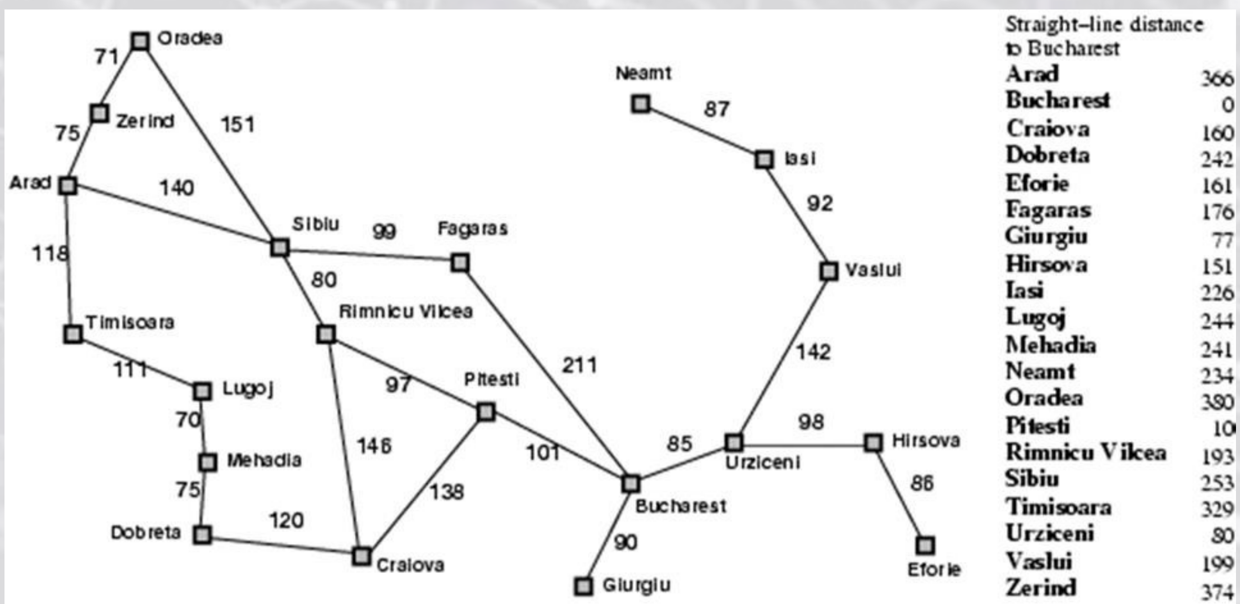

Oradea
71
Zerind
75
151
Arad
140
118
Sibiu
99
Fagaras
80
Rimnicu Vilcea
Timisoara
111
Lugoj
70
Mehadia
75
Dobreta
120
Craiova
146
97
Pitesti
138
101
211
Bucharest
90
Giurgiu
85
Urziceni
98
Hirsova
86
Eforie
142
Vaslui
92
Iasi
87
Neamt
Straight–line distance
to Bucharest
Arad 366
Bucharest 0
Craiova 160
Dobreta 242
Eforie 161
Fagaras 176
Giurgiu 77
Hirsova 151
Iasi 226
Lugoj 244
Mehadia 241
Neamt 234
Oradea 380
Pitesti 10
Rimnicu Vilcea 193
Sibiu 253
Timisoara 329
Urziceni 80
Vaslui 199
Zerind 374

A* search

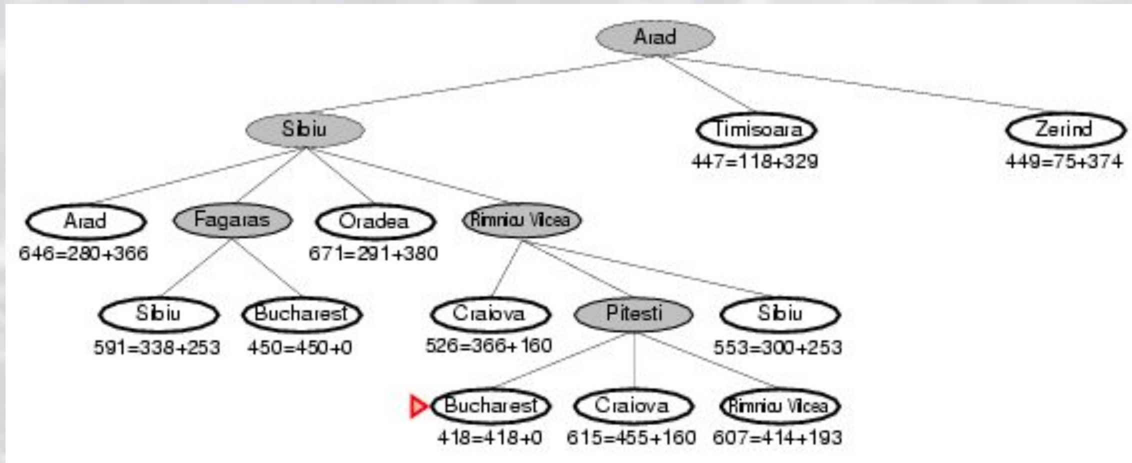
Arad
Sibiu
Timisoara
447=118+329
Zerind
449=75+374
Arad
646=280+366
Fagaras
Oradea
671=291+380
Rimnicu Vilcea
Sibiu
591=338+253
Bucharest
450=450+0
Craiova
526=366+160
Pitesti
Sibiu
553=300+253
Bucharest
418=418+0
Craiova
615=455+160
Rimnicu Vilcea
607=414+193

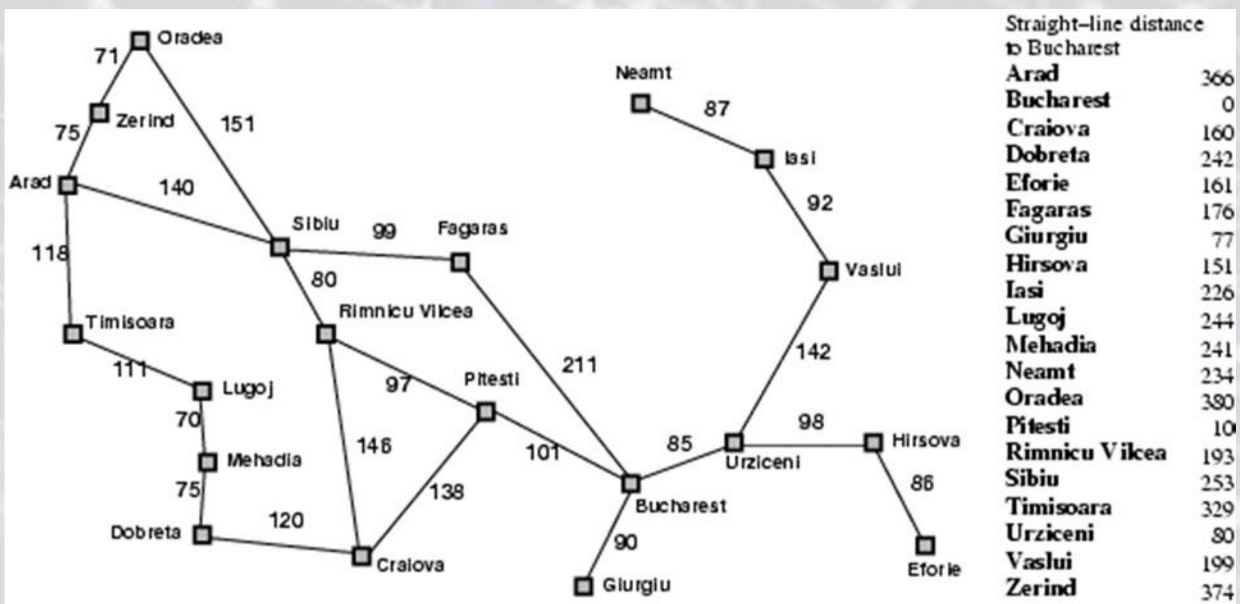
Oradea
71
Zerind
75
151
Arad
140
118
Sibiu
99
Fagaras
80
Rimnicu Vilcea
Timisoara
111
Lugoj
70
Mehadia
75
Dobreta
120
97
Pitesti
146
138
Craiova
101
211
Bucharest
90
Giurgiu
85
Urziceni
98
Hirsova
86
Eforie
142
Vaslui
92
Iasi
87
Neamt
Straight–line distance
to Bucharest
Arad 366
Bucharest 0
Craiova 160
Dobreta 242
Eforie 161
Fagaras 176
Giurgiu 77
Hirsova 151
Iasi 226
Lugoj 244
Mehadia 241
Neamt 234
Oradea 380
Pitesti 10
Rimnicu Vilcea 193
Sibiu 253
Timisoara 329
Urziceni 80
Vaslui 199
Zerind 374

Admissible heuristics

- A heuristic $h(n)$ is admissible if for every node n, $h(n) \leq h^*(n)$, where $h^*(n)$ is the true cost to reach the goal state from n.
- An admissible heuristic never overestimates the cost to reach the goal, i.e., it is optimistic.
- Example: $h_{SLD}(n)$ (never overestimates the actual road distance)
- Theorem: If $h(n)$ is admissible, A* using TREE- SEARCH is optimal.

Optimality of A* (proof)

- Suppose **some suboptimal goal** G_2 has been generated and is in the frontier. Let n be an unexpanded node in the frontier such that n is on a shortest path to an optimal goal G.

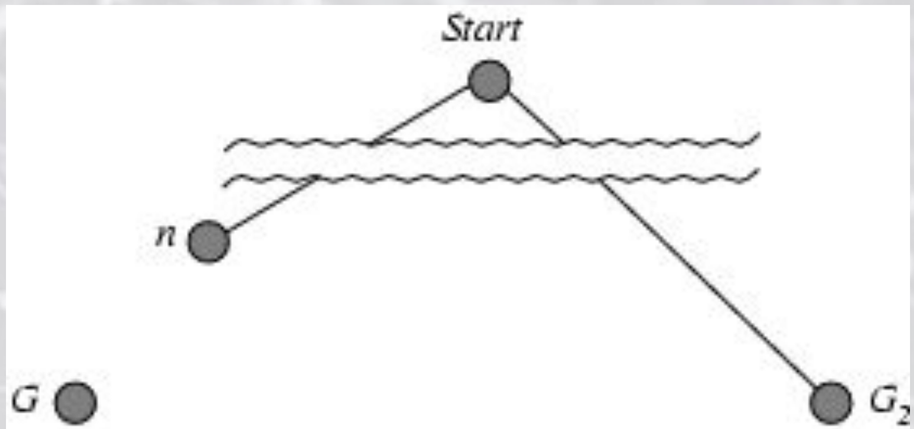

- $f(G_2) = g(G_2)$ since $h(G_2) = 0$
- $g(G_2) > g(G)$ since G_2 is suboptimal
- $f(G) = g(G)$ since $h(G) = 0$
- $f(G_2) > f(G)$ from above

Optimality of A* (proof)

- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G.

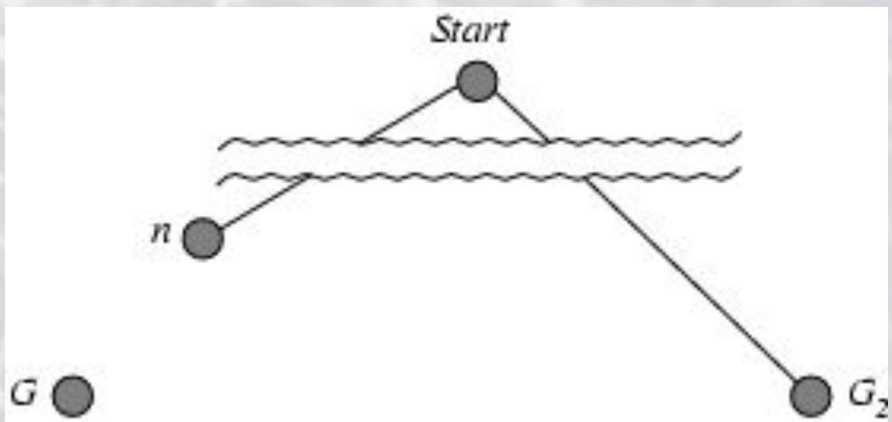

- $f(G_2) = g(G_2)$
- $g(G_2) > g(G)$
- $f(G) = g(G)$
- $f(G_2) > f(G)$

- $f(G_2) > f(G)$ (from above)
- $h(n) \leq h^*(n)$ (since h is **admissible**)

->$g(n) + h(n) \leq g(n) + h^*(n)$

- $f(n) \leq g(n) + h^*(n) < f(G) < f(G_2)$

Hence $f(G_2) > f(n)$, and A* will never select G_2 for expansion.

Consistent Heuristics

- A heuristic is **consistent** (or **monotonic**) if for every node *n*,
  every successor *n′* of *n* generated by any action *a:*

  $h(n) \le c(n,a,n') + h(n')$

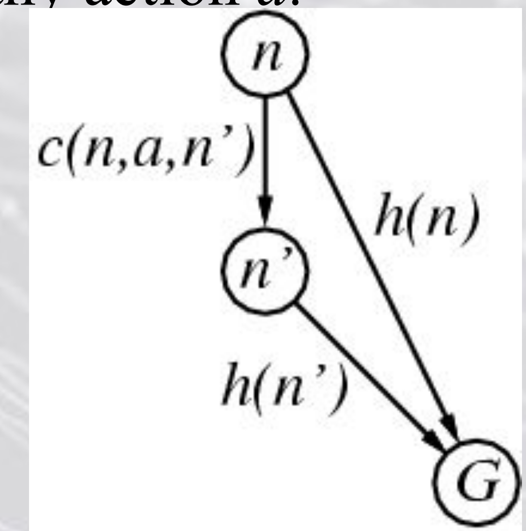

- If *h* is **consistent**, we have:

$$
\begin{aligned}
f(n') &= g(n') + h(n') \\
&= g(n) + c(n,a,n') + h(n') \\
&\geq g(n) + h(n) \\
&= f(n)
\end{aligned}
$$

i.e., *f(n)* is **non-decreasing along any path**.

Theorem: If *h(n)* is consistent, A* using GRAPH-SEARCH is optimal.

Optimality of A*

- A^* expands nodes in order of increasing f value.
- Gradually adds "f-contours" of nodes.
- Contour i has all nodes with $f=f_i$, where $f_i < f_{i+1}$.
- That is to say, nodes inside a given contour have f-costs less than or equal to contour value.

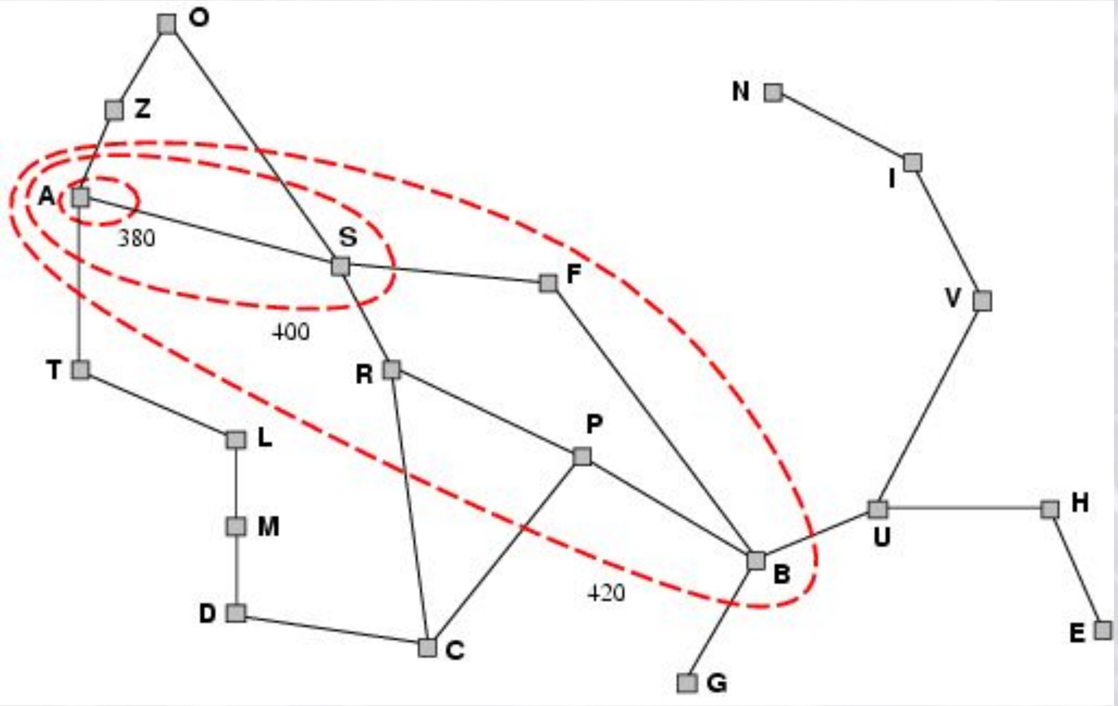

Properties of A*

- <u>Complete:</u> Yes (unless there are infinitely many nodes with $f \leq f(G)$).
- <u>Time:</u> Exponential.
- <u>Space:</u> Keeps all nodes in memory, so also exponential.
- <u>Optimal:</u> Yes (provided h admissible or consistent).
- <u>*Optimally Efficient*</u>: Yes (no algorithm with the same heuristic is guaranteed to expand fewer nodes).

- *NB*: Every consistent heuristic is also admissible (Pearl).

Q: What about the converse?

Admissible Heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance (i.e. 1-norm)

(i.e., no. of squares from desired location of each tile)

Q: Why are these admissible heuristics?

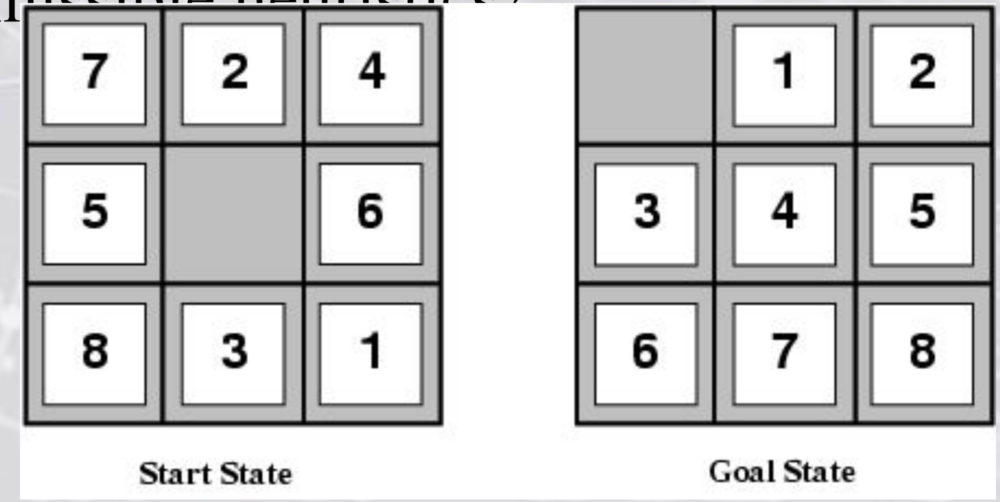

- $h_1(S)$ = ?
- $h_2(S)$ = ?

Admissible Heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance

(i.e., no. of squares from desired location of each tile)

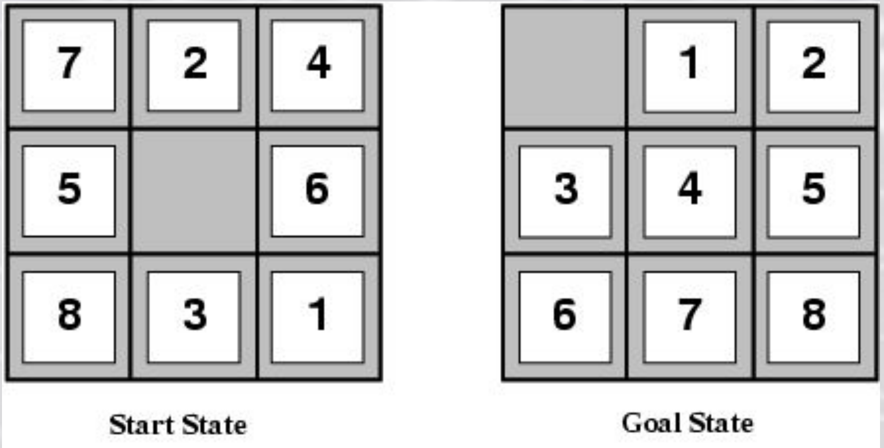

- $h_1(S) = ?$ 8
- $h_2(S) = ?$ 3+1+2+2+2+3+3+2 = 18

Dominance

- If $h_2(n) \geq h_1(n)$ for all n (both admissible), then h_2 dominates h_1.
- *Essentially, domination translates directly into efficiency: "h_2* is better for search.
- A* using h_2 will never expand more nodes than A* using h_1.
- Typical search costs (average number of nodes expanded):

d=12 IDS = 3,644,035 nodes
$A^*(h_1)$ = 227 nodes
$A^*(h_2)$ = 73 nodes

d=24 IDS = too many nodes
$A^*(h_1)$ = 39,135 nodes
$A^*(h_2)$ = 1,641 nodes

(IDS=iterative deepening

Memory Bounded Heuristic Search: Recursive BFS (best-first)

- How can we solve the memory problem for A* search?
- Idea: Try something like depth-first search, but let's <u>not forget everything about the branches we have partially explored</u>.
- *We remember the best f-value we have found so far in the branch we are deleting*.

Memory Bounded Heuristic Search: Recursive BFS

Best alternative over frontier nodes, which are not children:
i.e. do I want to back up?

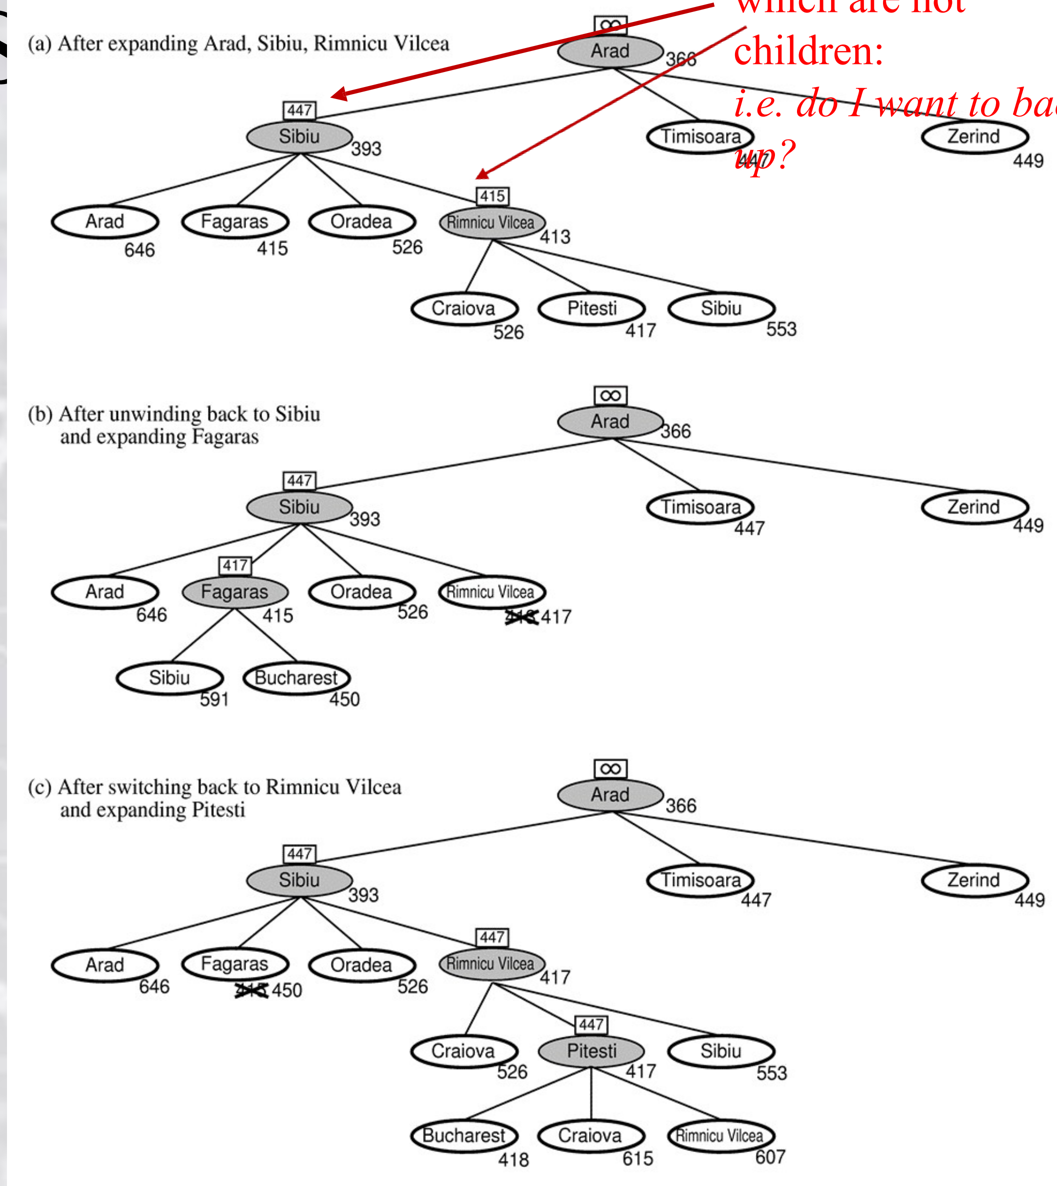

- RBFS changes its mind very often in practice. This is because f=g+h become more accurate (less optimistic) as we approach the goal. Hence, higher level nodes have smaller f-values and will be explored first.

- Problem: We should keep
- in memory whatever we can.

Simple Memory-Bounded A*

- This is like A*, but <u>when memory is full</u> we delete the worst node (largest f-value).
- Like RBFS, we remember the best descendent in the branch we delete.
- If there is a tie (equal f-values) we delete the oldest nodes first.
- Simple-MBA* finds the optimal *reachable* solution given the memory constraint (reachable means path from root to goal fits in memory).
- Can also use **iterative deepening** with A* (IDA*).
- Time can still be exponential.

Relaxed Problems

- A problem with fewer restrictions on the actions is called a relaxed problem.
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem. (why?)
- If the rules of the 8-puzzle are relaxed so that a tile can move anywhere, then $h_1(n)$ gives the shortest solution.
- If the rules are relaxed so that a tile can move to any adjacent square, then $h_2(n)$ gives the shortest solution.

Summary

- **Informed search** methods may have access to a **heuristic function** *h(n)* that estimates the cost of a solution from *n*.
- The generic **best-first search** algorithm selects a node for expansion according to an **evaluation function**.
- **Greedy best-first search** expands nodes with minimal h(n). It is not optimal, but is often efficient.
- **A*** search expands nodes with minimal f(n)=g(n)+h(n).
- A* s **complete** and **optimal**, provided that h(n) is admissible (for TREE-SEARCH) or consistent (for GRAPH-SEARCH).
- The space complexity of A* is still prohibitive.
- The performance of heuristic search algorithms depends